

Qubit control using reinforcement learning

Gabriele Borroni, UniMI, a.a. 2024/2025

Qubits

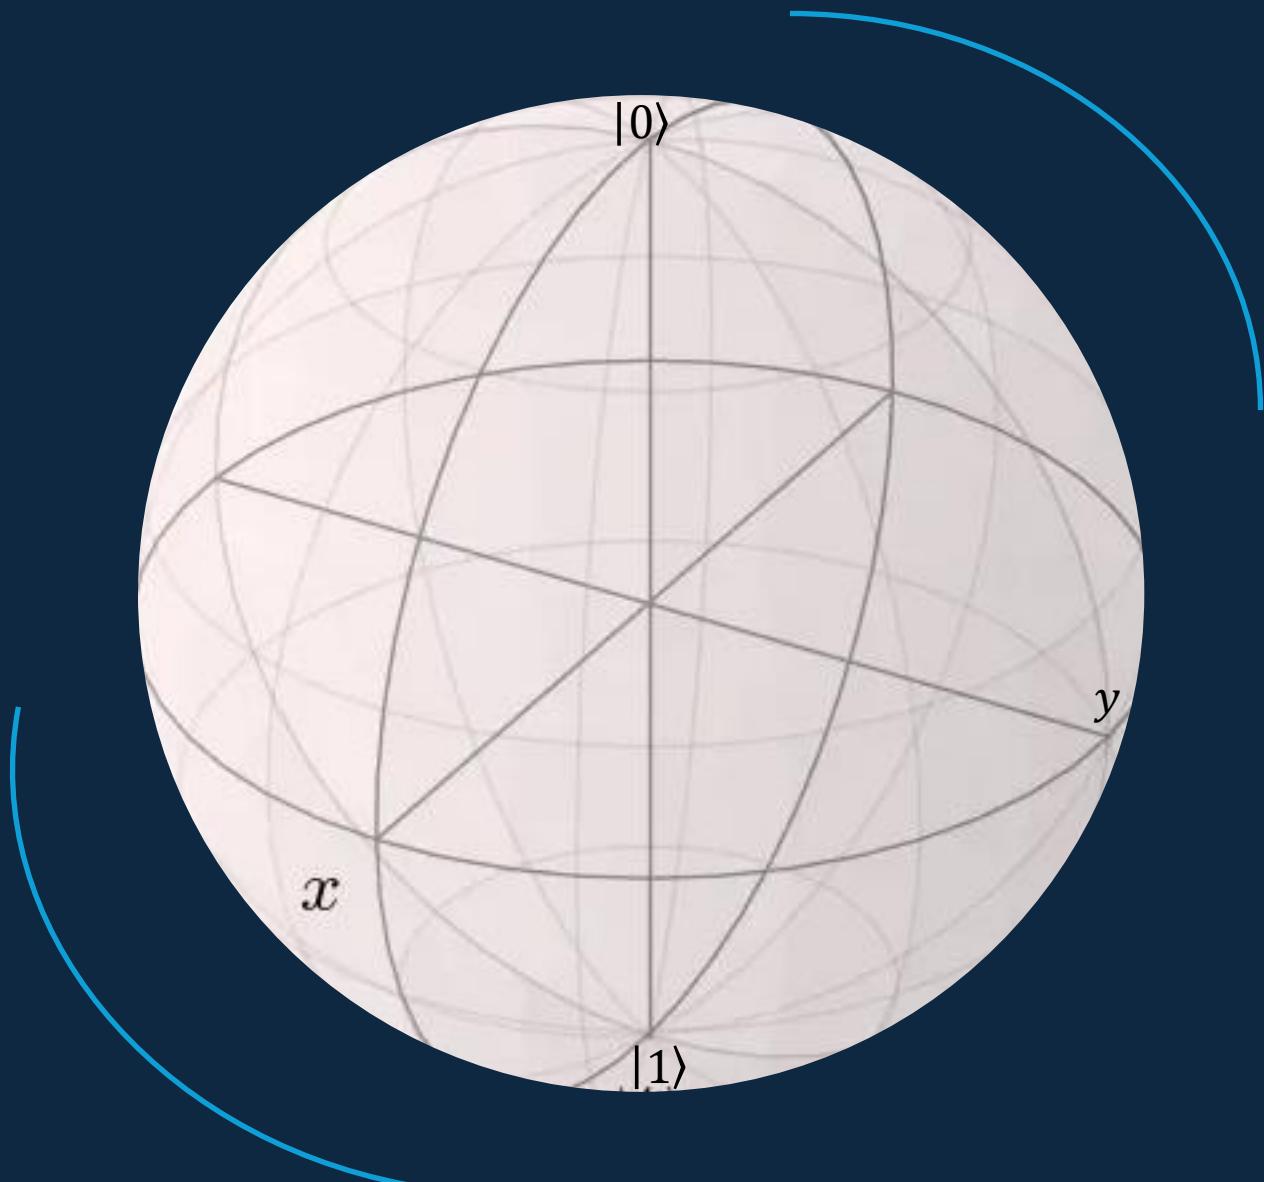
- 2-levels quantum systems:

$$\psi = \alpha|0\rangle + \beta|1\rangle$$

- Most effective representation is the Bloch-sphere:

$$\psi = \cos\left(\frac{\theta}{2}\right)|0\rangle + e^{i\phi} \sin\frac{\theta}{2} |1\rangle$$

$$(\sin\theta \cos\phi, \sin\theta \sin\phi, \cos\theta)$$



Spin Qubits

- Physical realizations of Qubit using an electron spin confined in a magnetic field: $\psi = \alpha|\uparrow\rangle + \beta|\downarrow\rangle$
- The Hamiltonian can be modeled as $H(t) = u_x(t)\sigma_x + u_y(t)\sigma_y$
- Where $U = (u_x, u_y)$ parametrizes the action of an extern E.M. field
- Time evolution:
$$\psi(t + \Delta t) = e^{-iH(t)\Delta t}|\psi(t)\rangle$$

```
import numpy as np
from scipy.linalg import expm #to esponentiate matrix

#define Pauli matrix
sigma_x = np.array([[0, 1], [1, 0]], dtype=np.complex128)
sigma_y = np.array([[0, -1j], [1j, 0]], dtype=np.complex128)
#define hamiltonian
def H(ux,uy):
    return ux*sigma_x+uy*sigma_y

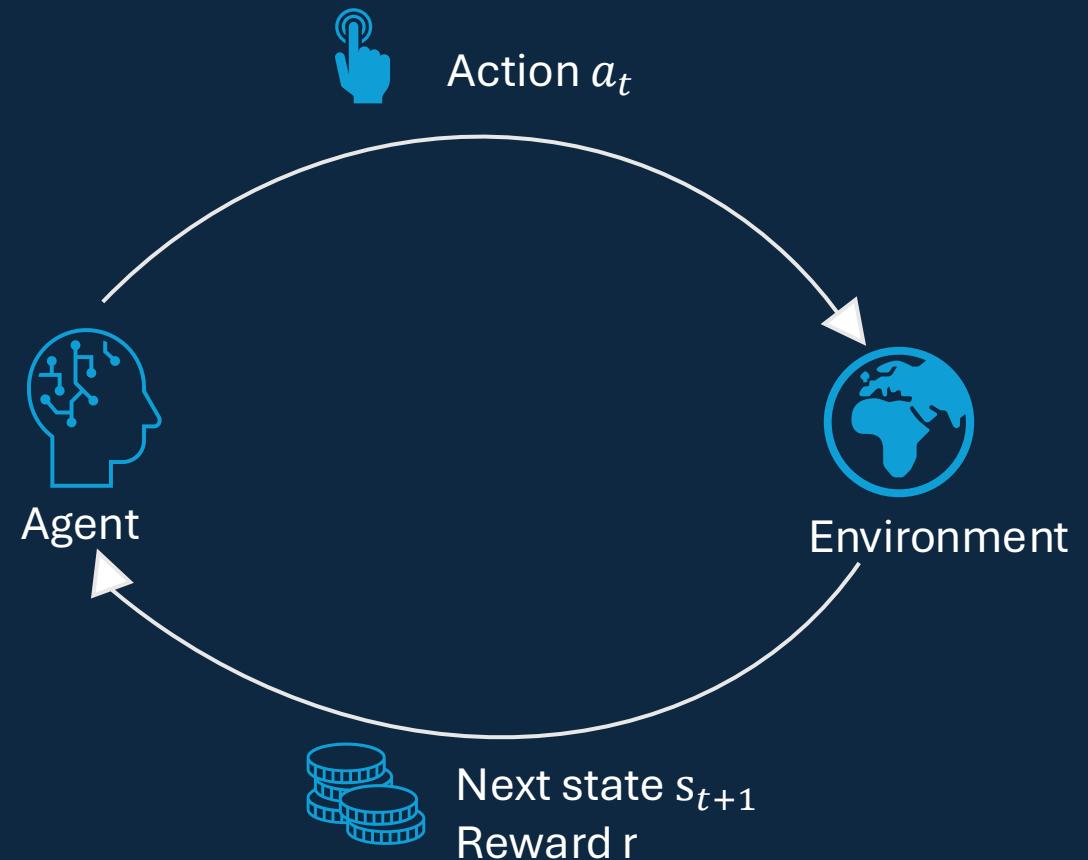
#define time evolution
def Evolve(state,ux,uy, deltat):
    U=expm(-1j*H(ux,uy)*deltat)
    new_state=U@state
    return new_state/np.linalg.norm(new_state)

#define fidelity (reward function) as the overlap |<psit|psi'||^2 (to be maximized)
def Fidelity(target_state,state):
    f= np.abs(np.vdot(target_state, state))**2
    return f
```

Reinforcement learning

- Agent interacts with the environment and receives a reward based on its action.
- The better the action, the greater the reward
- The agent follows a policy, which tells it the action to take in each state:
$$\pi(s_t) = a_t$$
- The agent aims to find a policy that maximizes the expected cumulative (discounted) reward:

$$\pi^* = \arg \max_{\pi} E \left[\sum_{t \geq 0} \gamma^t r_t \mid \pi \right]$$



Deep Q-Learning (DQN)

- The core is the action-value function:

$$Q^\pi(s, a) = E[\sum_{t=0}^{\infty} \gamma^t r_t \mid s_0 = s, a_0 = a, \pi]$$

- The best policy is the one that follows the optimal Q-function:

$$Q^* = \max_{\pi} Q^\pi(s, a) \rightarrow \pi^* = \arg \max_a Q^*(s, a)$$

- It can be recursively computed using the bellman equation:

$$Q^*(s, a) = \mathbb{E}[r + \gamma \max_{a'}(Q(s', a')) \mid s, a]$$

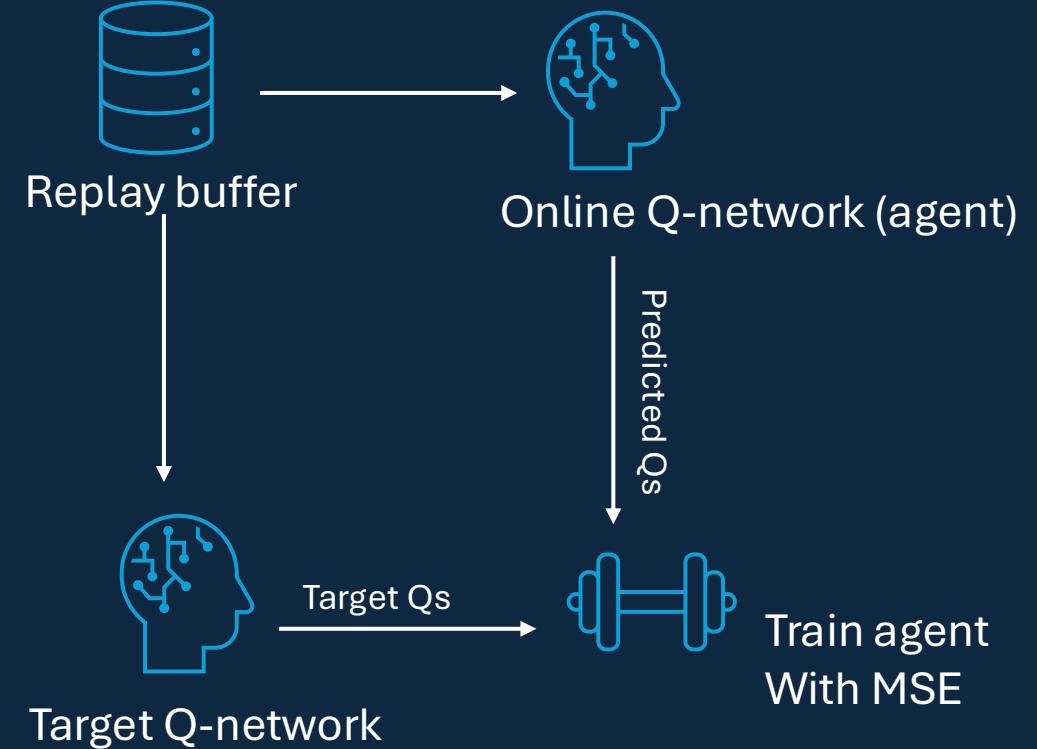
- At each step, the update follows:

$$Q^{i+1}(s, a) = \mathbb{E}[r + \gamma \max_{a'}(Q^i(s', a')) \mid s, a] \rightarrow Q^*$$

Training loop

- $Q^\pi(s, a) = E[\sum_{t=0}^{\infty} \gamma^t r_t \mid s_0 = s, a_0 = a, \pi]$
evaluates expected reward of a state-action pair, following the current policy. Here it is parametrized by a DNN.
- the agent stores experiences in a replay buffer, then samples minibatches to compute target values y using a target DNN, (a copy of the agent updated periodically)
$$y = r + \gamma \max_{a'} Q(s', a'; \theta^-)$$
- The online network is trained by minimizing the MSE loss between target and predicted Q-values:

$$L(\theta) = \text{MSE}(Q(s_i, a_i, \theta), y_i)$$



First approach: discretized action space

```
def build_agent(obs_shape, actions):
    model=tf.keras.models.Sequential([
        tf.keras.layers.Flatten(input_shape=obs_shape),
        tf.keras.layers.Dense(256, activation='relu'),
        tf.keras.layers.Dense(256, activation='relu'),
        tf.keras.layers.Dense(actions, activation='linear')])
    return model

#now define policy (it is a guideline on how the agent will behave during training). i choose
#epsilon-greedy, which performs a random action with probability epsilon, otherwise it performs
#the action which maximises the expected reward Q(s,a)
def epsilon_greedy_policy(state, actions, agent, epsilon=0.1):
    if(np.random.rand()<=epsilon):
        return np.random.choice(actions)#random action with Prob=epsilon
    else:
        q_values=agent.predict(state[np.newaxis], verbose=0)
        return np.argmax(q_values)
```

```
class Qbitenvironment_discrete(gym.Env):

    def __init__(self):
        super().__init__()
        #define possible actions: discrete tuning of fields
        self.u_vals=np.linspace(-1.0, 1.0, 5)
        self.actions = [(ux, uy) for ux in self.u_vals for uy in self.u_vals]
        self.action_space = spaces.Discrete(len(self.actions)) #discrete space for now
        # (ax), (ay), (az)
        self.observation_space = spaces.Box(low=-1.0, high=1.0, shape=(3,), dtype=np.float32)
        #timestep
        self.dt = 0.15
        self.sigma_x = np.array([[0, 1], [1, 0]], dtype=np.complex128)
        self.sigma_y = np.array([[0, -1j], [1j, 0]], dtype=np.complex128)
        self.sigma_z = np.array([[1, 0], [0, -1]], dtype=np.complex128)
        self.psi_target = np.array([0.0, 1.0], dtype=np.complex128) # es. |+
        self.max_steps = 50
        self.reset()

    def reset(self, seed=None, options=None):
        #psi = np.random.randn(2) + 1j*np.random.randn(2) #random initial state
        #psi=np.linalg.norm(psi)
        #self.psi=psi
        self.psi = np.array([1.0, 0.0], dtype=np.complex128) # |0>

        obs = self._get_observation()
        self.current_step = 0
        return obs, {}

    def step(self, action_idx):
        ux, uy = self.actions[action_idx]
        self.psi = Evolve(self.psi, ux, uy, self.dt)
        obs = self._get_observation()
        reward = Fidelity(self.psi_target, self.psi)
        self.current_step += 1
        truncated=False
        terminated=False

        if reward > 0.95:
            terminated=True

        if (self.current_step >= self.max_steps):
            truncated = True
        return obs, reward, terminated, truncated, {}

    def _get_observation(self):
        sx = np.vdot(self.psi, self.sigma_x @ self.psi).real
        sy = np.vdot(self.psi, self.sigma_y @ self.psi).real
        sz = np.vdot(self.psi, self.sigma_z @ self.psi).real
        return np.array([sx, sy, sz], dtype=np.float32)
```

Training loop

```
env = Qbitenvironment_discrete()
obs_shape = env.observation_space.shape
actions = env.action_space.n

agent = build_agent(obs_shape, actions) # online
target_agent = build_agent(obs_shape, actions) # target
agent.compile(optimizer=tf.keras.optimizers.Adam(learning_rate=1e-3), loss='mse')
target_agent.set_weights(agent.get_weights())

episodes = 1000
update_target_every = 3
memory = deque(maxlen=20000)
batch_size = 10
gamma = 0.99

reward_history = []

for episode in range(episodes):
    state, info = env.reset()
    episode_reward = 0.0
    terminated = False
    truncated = False

    while not (terminated or truncated):
        action = epsilon_greedy_policy(state, actions, agent)
        next_state, reward, terminated, truncated, _ = env.step(action)
        memory.append((state, action, reward, next_state, terminated, truncated))
        episode_reward += reward
        state = next_state

    if len(memory) >= batch_size:
        minibatch = random.sample(memory, batch_size)
        states_mb = np.stack([mb[0] for mb in minibatch])
        actions_mb = np.array([mb[1] for mb in minibatch])
        rewards_mb = np.array([mb[2] for mb in minibatch])
        next_states_mb = np.stack([mb[3] for mb in minibatch])
        terminated_mb = np.array([mb[4] for mb in minibatch], dtype=bool)
        truncated_mb = np.array([mb[5] for mb in minibatch], dtype=bool)

        # current Q-values (online)
        q_current = agent.predict(states_mb, verbose=0) # (B, A)

        # predicted Qs for next states (target)
        next_q_all = target_agent.predict(next_states_mb, verbose=0) # (B, A)
        max_next_q = next_q_all.max(axis=1) # (B,)
        done = (terminated_mb | truncated_mb).astype(np.float32)

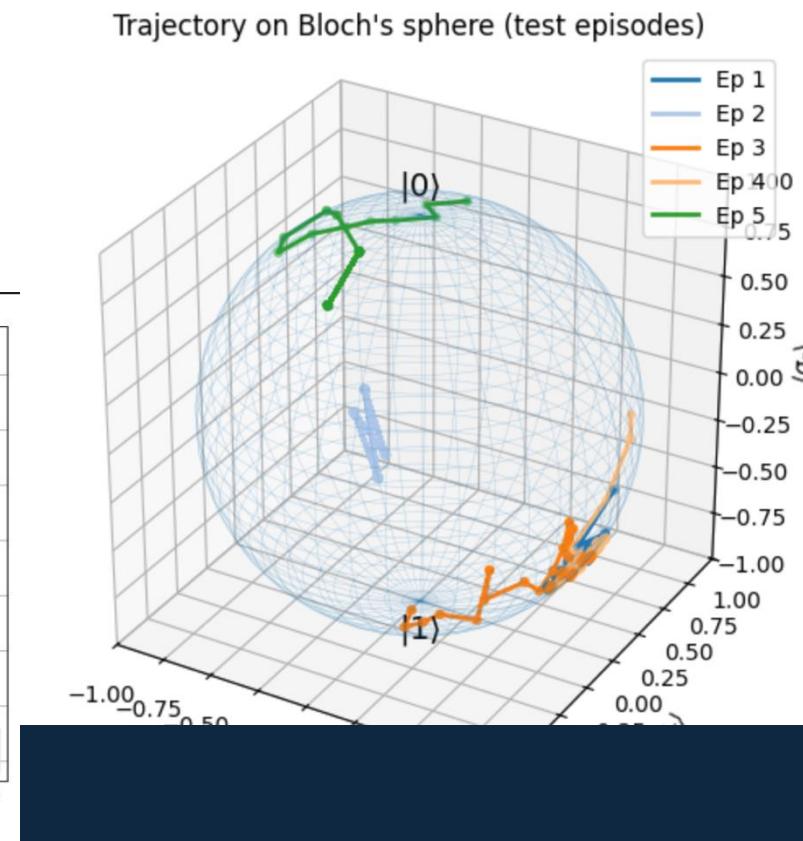
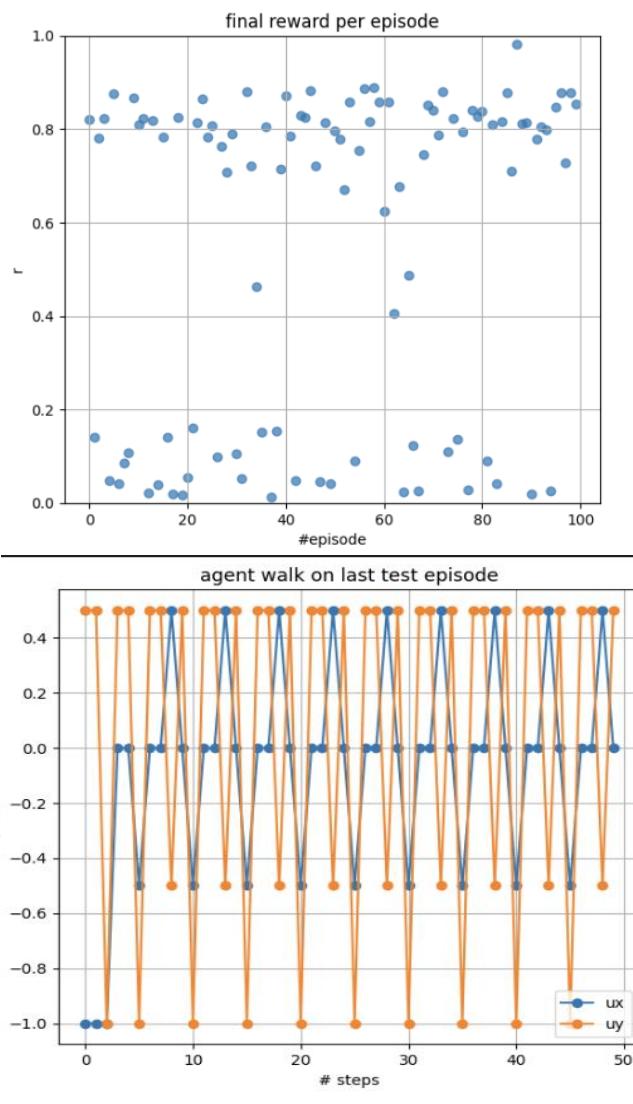
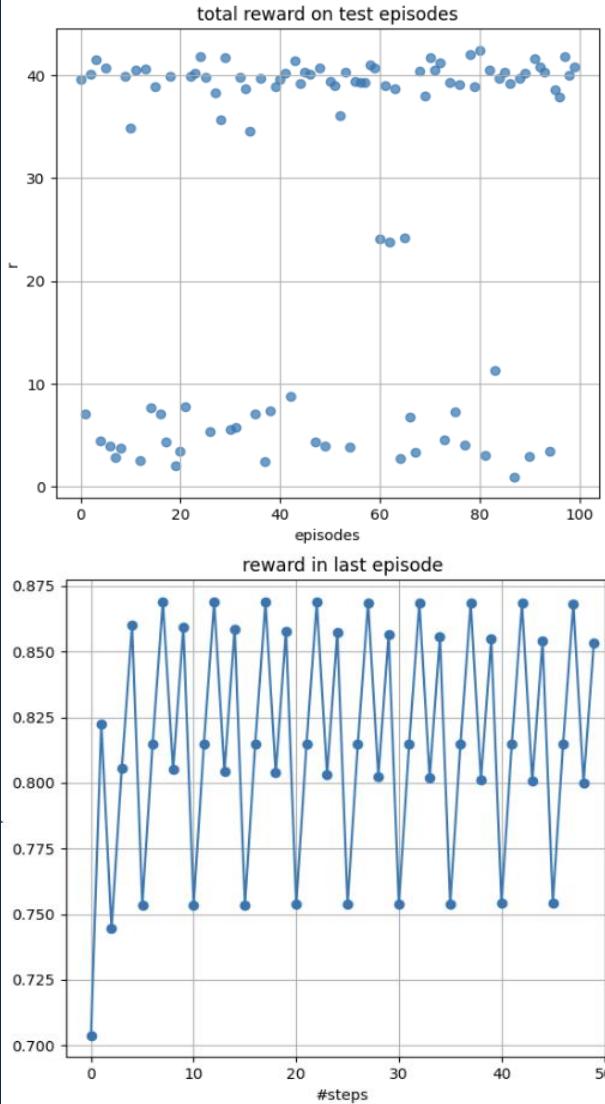
        # Bellman targets (reward is r if episode is finished, else it's predicted by best Q target)
        q_target_vec = rewards_mb + gamma * max_next_q * (1.0 - done)
        # update the q prediction for the selected action only
        targets = q_current.copy()
        targets[np.arange(batch_size), actions_mb] = q_target_vec
        # train online network
        agent.fit(states_mb, targets, epochs=1, verbose=0)

    # after episode's end
    if (episode + 1) % update_target_every == 0:
        target_agent.set_weights(agent.get_weights())

    reward_history.append(episode_reward)
    if episode >= 10:
        print(f"Ep {episode:4d} | AvgR(10)={np.mean(reward_history[-10:]):.3f}")
    else:
        print(f"Ep {episode:4d} | R={episode_reward:.3f}")

    env.close()
```

Results



Continuous action space- Soft Actor Critic (SAC)

Ingredients:

- Actor learns to maximize:

$$J(\pi) = \mathbb{E} \left[\sum_t \gamma^t r(s, t) - \alpha \log(\pi(a|s)) \right]$$

- Stochastic policy (actor): π_θ .

- Actor returns the parameters of a PDF (Gaussian) from which the action is sampled

- 2-Q-values networks (critic)

- Temperature for entropy regularization α

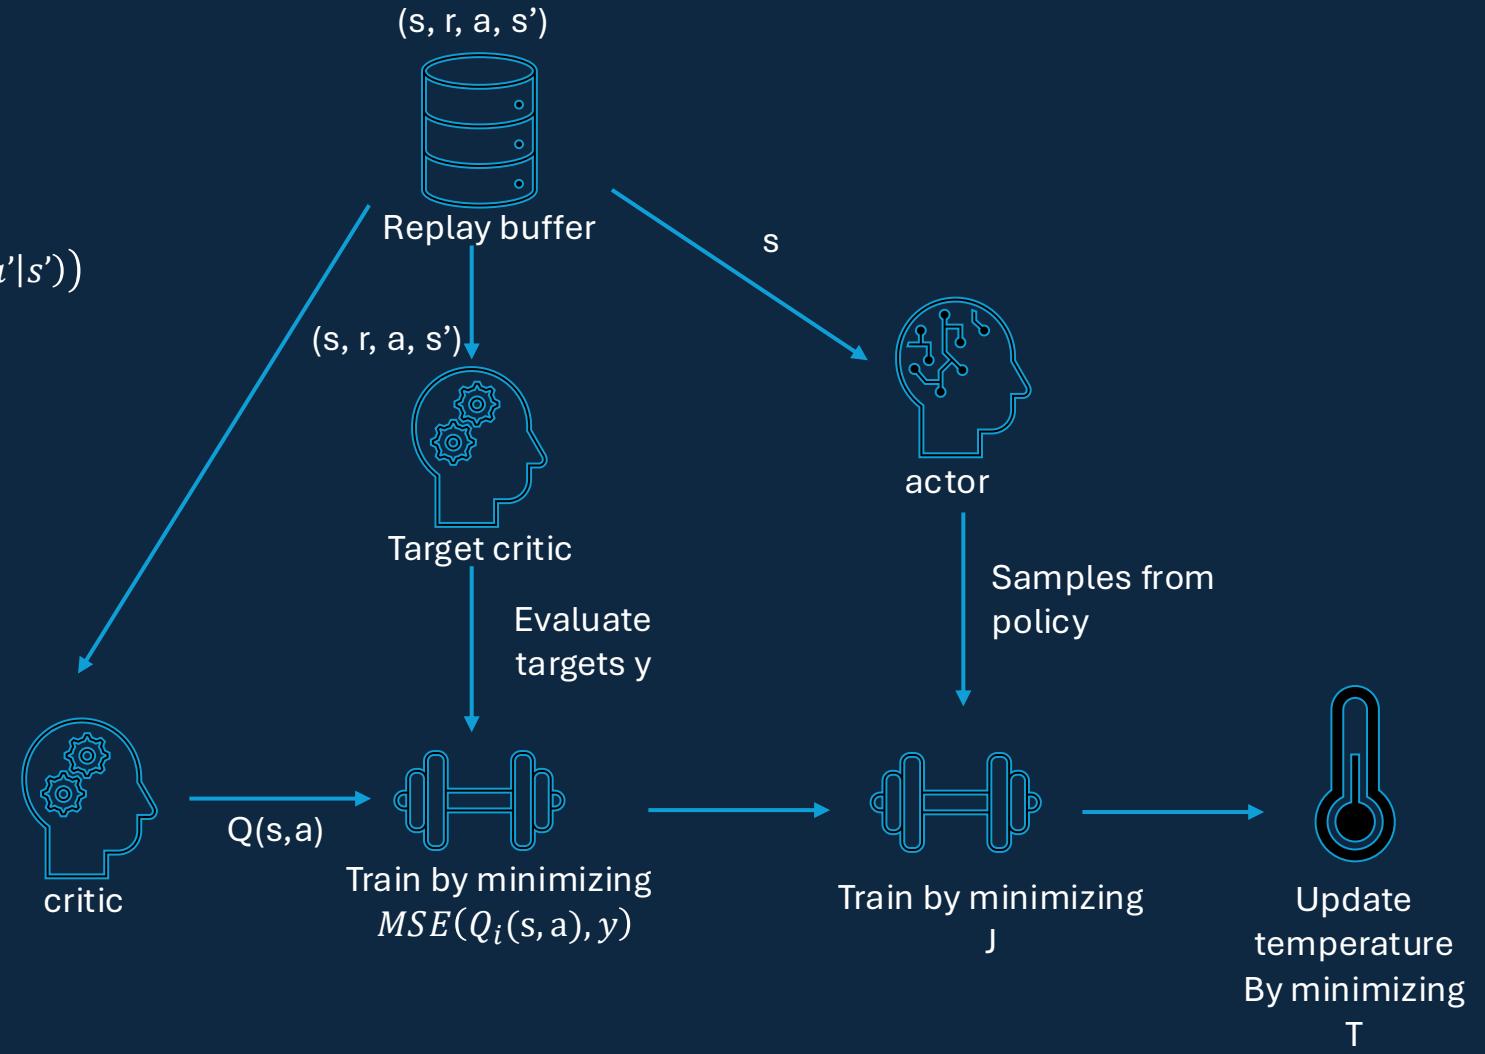
- Replay buffer

Advantages:

- Handles continuous action space efficiently thanks to the stochastic policy
- Stable exploration-exploitation trade off through entropy regularization
- Reduced overestimation bias by using two Q-networks

Training loop

- Target evaluation:
$$y_i = r + \gamma (\min(Q_1(s', a'), Q_2(s', a')) - \alpha \log \pi_\theta(a'|s'))$$
- Actor loss:
$$J(\pi) = \mathbb{E}[\alpha \log \pi - \min(Q_1(s, a), Q_2(s, a))]$$
- Temperature update loss:
$$T(\alpha) = \mathbb{E}[-\alpha(\log(\pi(s|a) + H_{target})]$$



```

env = Qbitenvironment_continuum()
obs_shape = env.observation_space.shape[0]
actions_shape = env.action_space.shape[0]
#define Q(s,a) networks
def build_Q(obs_shape, actions_shape):
    input_obs= tf.keras.layers.Input(shape=(obs_shape,))
    input_actions= tf.keras.layers.Input(shape=(actions_shape,))
    input=tf.keras.layers.concatenate([input_obs, input_actions])
    h1=tf.keras.layers.Dense(256, activation='relu')(input)
    h2=tf.keras.layers.Dense(256, activation='relu')(h1)
    Q=tf.keras.layers.Dense(1, activation='linear')(h2)
    model=tf.keras.Model(inputs=[input_obs, input_actions], outputs=Q)
    return model

#define policy network
def build_policy(obs_shape, actions_shape):
    inputs= tf.keras.layers.Input(shape=(obs_shape,))
    h1= tf.keras.layers.Dense(256, activation='relu')(inputs)
    h2=tf.keras.layers.Dense(256, activation='relu')(h1)
    mu=tf.keras.layers.Dense(actions_shape, activation='linear')(h2)#first output layer
    logsigma=tf.keras.layers.Dense(actions_shape, activation='linear')(h2)#second output layer
    model=tf.keras.Model(inputs=inputs, outputs=[mu,logsigma])
    return model

def sample_from_policy(mu, log_sigma):
    sigma = tf.exp(log_sigma)
    eps = tf.random.normal(shape=tf.shape(mu))
    a_prime = mu + sigma * eps
    a = tf.tanh(a_prime)
    #logN(a)=logN(a'|mu,o2)-iΣlog(1-tanh2(ai')) (change of variables)
    # log N(a'|mu, sigma^2)
    log_prob_gauss = -0.5 * (
        tf.square((a_prime - mu) / (sigma + 1e-8)) +
        2.0 * log_sigma +
        tf.math.log(2.0 * np.pi)
    )
    log_prob_gauss = tf.reduce_sum(log_prob_gauss, axis=-1, keepdims=True)
    log_det_jac = tf.reduce_sum(
        tf.math.log(1.0 - tf.square(tf.tanh(a_prime)) + 1e-6),
        axis=-1, keepdims=True
    )

    log_pi = log_prob_gauss - log_det_jac
    return a, log_pi

```

```

def Train_Qs(Q1, Q2,state_mb, action_mb, targets):
    with tf.GradientTape(persistent=True) as tape:
        #calculate predictions
        predicted_q1s = Q1([state_mb, action_mb], training=True)
        predicted_q2s = Q2([state_mb, action_mb], training=True)
        #evaluate mse
        mse1= tf.reduce_mean(tf.square(predicted_q1s-targets))
        mse2= tf.reduce_mean(tf.square(predicted_q2s-targets))

        #evaluate gradients
        grad1= tape.gradient(mse1, Q1.trainable_variables)
        grad2= tape.gradient(mse2, Q2.trainable_variables)
        #perform one optimization step using adam
        optimizer_q1.apply_gradients(zip(grad1, Q1.trainable_variables))#apply_gradients
        optimizer_q2.apply_gradients(zip(grad2, Q2.trainable_variables))
    del tape
    return mse1, mse2

def Train_policy(state_mb, Q1, Q2, policy, log_alpha):
    with tf.GradientTape() as tape:
        mu, logsigma = policy(state_mb, training=True)
        #evaluate action and logpi
        action, log_pi = sample_from_policy(mu, logsigma)
        #update temperature
        entropy = Update_alpha( log_alpha, log_pi)
        #evaluate best estimation of q for the sampled action
        q1=Q1([state_mb, action], training=False)
        q2=Q2([state_mb, action], training=False)
        q=tf.minimum(q1,q2)
        #evaluate loss
        loss=-(tf.reduce_mean(q-tf.exp(log_alpha)*log_pi))#negative loss, i want to maximize -loss
        #evaluate gradient
        grad=tape.gradient(loss, policy.trainable_variables)
        #perform optimization step
        optimizer_pi.apply_gradients(zip(grad, policy.trainable_variables))
    return loss, entropy

def update_target_network(Q, target_Q, rho=0.995):
    #rho is a parameter that handles the "velocity" of updating the target Qs with the main Qs parameters.
    #e.g. rho=0-> copy all the parameters of Q into the target network.
    #by doing so, the target network weights become a mobile average of main network's (stabilize bootstrap of Qs
    main_weights=Q.get_weights()
    target_weights=target_Q.get_weights()
    updated_target_weights= [rho * tw + (1 - rho) * mw for mw, tw in zip(main_weights, target_weights)]
    target_Q.set_weights(updated_target_weights)

def Update_alpha(log_alpha, logpi, Htarget=-2): #from the paper: Htarget=-dim(action space)
    with tf.GradientTape() as tape:
        alpha=tf.exp(log_alpha)
        loss= tf.reduce_mean(alpha*(logpi+Htarget))
        grad =tape.gradient(loss, [log_alpha])
        optimizer_alpha.apply_gradients(zip(grad, [log_alpha]))
    return loss

```

Training loop

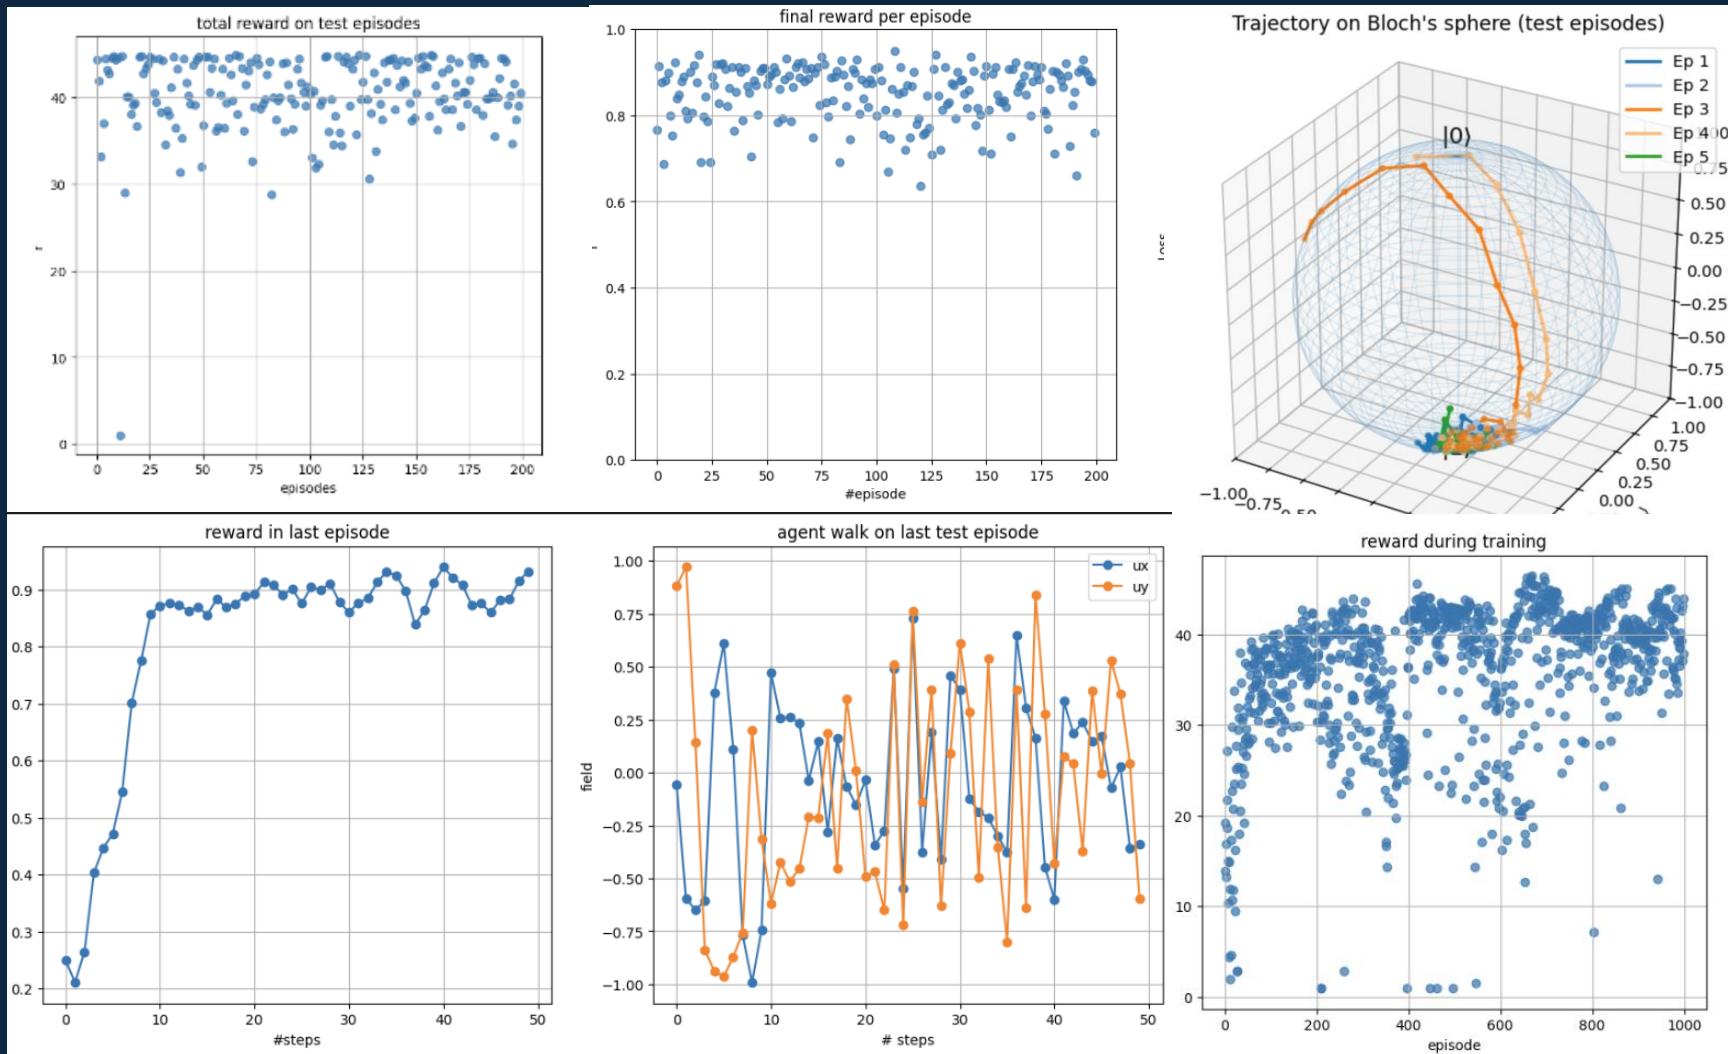
```
for episode in range(episodes):
    state, info= env.reset()
    terminated= False
    truncated = False
    episode_reward=0.0
    steps=0
    while not (terminated or truncated):
        mu, logsigma=policy(state[np.newaxis], training=False)
        action, _=sample_from_policy(mu,logsigma)
        next_state, reward, terminated, truncated, _=env.step(action.numpy()[0])
        done = terminated or truncated
        replay_buffer.append((state, action, reward, next_state, done) )
        state=next_state
        episode_reward+=reward
        steps+=1
        #training
        if(len(replay_buffer)>=batch_size):
            #random sample a batch of transitions from replay buffer
            minibatch= random.sample(replay_buffer, batch_size)
            #targets yi:
            state_mb, action_mb, reward_mb, next_state_mb, done_mb=map(np.array, zip(*minibatch))
            reward_mb = tf.reshape(tf.convert_to_tensor(reward_mb, dtype=tf.float32), (-1,1))
            done_mb = tf.reshape(tf.convert_to_tensor(done_mb, dtype=tf.float32), (-1,1))
            state_mb = tf.convert_to_tensor(state_mb, dtype=tf.float32)
            action_mb = tf.convert_to_tensor(action_mb, dtype=tf.float32)
            action_mb = tf.squeeze(action_mb, axis=1) # (64,1,2) → (64,2)

            next_state_mb = tf.convert_to_tensor(next_state_mb, dtype=tf.float32)

            #evaluate target Qs predictions and take the min
            mutrial, logsigmatrial = policy(next_state_mb, training= False)
            trial_action, log_pi = sample_from_policy(mutrial, logsigmatrial)
            target_qs1=target_Q1([next_state_mb, trial_action], training= False)
            target_qs2=target_Q2([next_state_mb, trial_action], training=False)
            #evaluate targets
            y=reward_mb+gamma*(1-done_mb)*(tf.minimum(target_qs1, target_qs2)- tf.exp(log_alpha)*log_pi)
            #update Q networks
            mse1, mse2= Train_Qs(Q1, Q2, state_mb, action_mb, y)
            critic_losses1.append(mse1)
            critic_losses2.append(mse2)
            #update policy network (and temperature)
            actor_loss, entropy_loss=Train_policy(state_mb, Q1, Q2, policy, log_alpha)
            actor_losses.append(actor_loss)
            entropy_losses.append(entropy_loss)
            #update target networks
            if(steps)%update_target_every==0:
                update_target_network(Q1, target_Q1)
                update_target_network(Q2, target_Q2)

            reward_history.append(episode_reward)
            if episode >= 10:
                print(f"Ep {episode:4d} | R={episode_reward:.3f}| AvgR(10)={np.mean(reward_history[-10:]):.3f}| steps=(steps)")
            else:
                print(f"Ep {episode:4d} | R={episode_reward:.3f}")
```

Results



Possible improvements

- Extend the system to more than one Qbit
- Extend to different Hamiltonians
- Introduce noise models

Citations

- DQN: "Playing Atari with deep reinforcement learning"- DeepMind- <https://arxiv.org/pdf/1312.5602.pdf>
- Soft Actor-Critic: Off-Policy Maximum Entropy Deep Reinforcement Learning with a Stochastic Actor- Tuomas Haarnoja, Aurick Zhou, Pieter Abbeel, Sergey Levine. <https://arxiv.org/abs/1801.01290>
- "Modern" version of the SAC algorithm by openAI: <https://spinningup.openai.com/en/latest/algorithms/sac.html#soft-actor-critic>
- Furtherly improved version of the SAC algorithm, including automatic temperature regularization: Soft Actor-Critic Algorithms and Applications- T.Haarnoja et al. <https://arxiv.org/abs/1801.01290>
- Gymnasium custom environments: https://gymnasium.farama.org/introduction/create_custom_env/